

# Ch4\_Lab

Adrien Osakwe

## Stock Market Data

Working with the *Smarket* data set in the ISLR2 package. Looking to use data from 1250 days of stock indices to predict stock price direction.

```
library(ISLR2)
```

```
## Warning: package 'ISLR2' was built under R version 4.1.2
```

```
names(Smarket)
```

```
## [1] "Year"      "Lag1"      "Lag2"      "Lag3"      "Lag4"      "Lag5"
## [7] "Volume"    "Today"     "Direction"
```

```
dim(Smarket)
```

```
## [1] 1250    9
```


```
summary(Smarket)
```

```
##      Year      Lag1      Lag2      Lag3
## Min.   :2001   Min.   :-4.922000   Min.   :-4.922000   Min.   :-4.922000
## 1st Qu.:2002   1st Qu.: -0.639500   1st Qu.: -0.639500   1st Qu.: -0.640000
## Median :2003   Median : 0.039000   Median : 0.039000   Median : 0.038500
## Mean   :2003   Mean   : 0.003834   Mean   : 0.003919   Mean   : 0.001716
## 3rd Qu.:2004   3rd Qu.: 0.596750   3rd Qu.: 0.596750   3rd Qu.: 0.596750
## Max.   :2005   Max.   : 5.733000   Max.   : 5.733000   Max.   : 5.733000
##      Lag4      Lag5      Volume      Today
## Min.   :-4.922000   Min.   :-4.92200   Min.   :0.3561   Min.   :-4.922000
## 1st Qu.: -0.640000   1st Qu.: -0.64000   1st Qu.:1.2574   1st Qu.: -0.639500
## Median : 0.038500   Median : 0.03850   Median :1.4229   Median : 0.038500
## Mean   : 0.001636   Mean   : 0.00561   Mean   :1.4783   Mean   : 0.003138
## 3rd Qu.: 0.596750   3rd Qu.: 0.59700   3rd Qu.:1.6417   3rd Qu.: 0.596750
## Max.   : 5.733000   Max.   : 5.73300   Max.   :3.1525   Max.   : 5.733000
## Direction
## Down:602
## Up :648
##
##
##
##
```

```
## Generate correlation matrix using the quantitative variables
cor(Smarket[, -9])
```

```
##           Year      Lag1      Lag2      Lag3      Lag4
## Year  1.00000000  0.029699649  0.030596422  0.033194581  0.035688718
## Lag1  0.02969965  1.000000000 -0.026294328 -0.010803402 -0.002985911
## Lag2  0.03059642 -0.026294328  1.000000000 -0.025896670 -0.010853533
## Lag3  0.03319458 -0.010803402 -0.025896670  1.000000000 -0.024051036
## Lag4  0.03568872 -0.002985911 -0.010853533 -0.024051036  1.000000000
## Lag5  0.02978799 -0.005674606 -0.003557949 -0.018808338 -0.027083641
## Volume 0.53900647  0.040909908 -0.043383215 -0.041823686 -0.048414246
## Today 0.03009523 -0.026155045 -0.010250033 -0.002447647 -0.006899527
##           Lag5      Volume      Today
## Year  0.029787995  0.53900647  0.030095229
## Lag1 -0.005674606  0.04090991 -0.026155045
## Lag2 -0.003557949 -0.04338321 -0.010250033
## Lag3 -0.018808338 -0.04182369 -0.002447647
## Lag4 -0.027083641 -0.04841425 -0.006899527
## Lag5  1.000000000 -0.02200231 -0.034860083
## Volume -0.022002315  1.00000000  0.014591823
## Today -0.034860083  0.01459182  1.000000000
```

```
# A quick scan of the matrix indicates a correlation between the year and the
#Volume
attach(Smarket)
boxplot(Volume ~ Year)
```



## Logistic Regression

Using log regression to predict direction with lag[1:5] and volume

```
#Use glm() with family = binomial to cal log-reg
glm.fits <- glm(Direction ~ Lag1 + Lag2 + Lag3 + Lag4 + Lag5 + Volume, data =
               Smarket, family = binomial)

summary(glm.fits)
```

```
##
## Call:
## glm(formula = Direction ~ Lag1 + Lag2 + Lag3 + Lag4 + Lag5 +
##       Volume, family = binomial, data = Smarket)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.446  -1.203   1.065   1.145   1.326
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept) -0.126000   0.240736  -0.523   0.601
## Lag1        -0.073074   0.050167  -1.457   0.145
## Lag2        -0.042301   0.050086  -0.845   0.398
```

```
## Lag3      0.011085  0.049939  0.222    0.824
## Lag4      0.009359  0.049974  0.187    0.851
## Lag5      0.010313  0.049511  0.208    0.835
## Volume    0.135441  0.158360  0.855    0.392
##
## (Dispersion parameter for binomial family taken to be 1)
##
## Null deviance: 1731.2  on 1249  degrees of freedom
## Residual deviance: 1727.6  on 1243  degrees of freedom
## AIC: 1741.6
##
## Number of Fisher Scoring iterations: 3
```

```
#smallest p-value is for lag1. as b1 is negative, a positive return lag1 would
#predict and positive return today
summary(glm.fits)$coef
```

```
##           Estimate Std. Error   z value Pr(>|z|)
## (Intercept) -0.126000257 0.24073574 -0.5233966 0.6006983
## Lag1        -0.073073746 0.05016739 -1.4565986 0.1452272
## Lag2        -0.042301344 0.05008605 -0.8445733 0.3983491
## Lag3         0.011085108 0.04993854  0.2219750 0.8243333
## Lag4         0.009358938 0.04997413  0.1872757 0.8514445
## Lag5         0.010313068 0.04951146  0.2082966 0.8349974
## Volume       0.135440659 0.15835970  0.8552723 0.3924004
```

```
## Running predict() here requires us to set the type parameter to "response" to
# output a posterior probability
contrasts(Direction) #this informs us that the probability is for today being Up
```

```
##      Up
## Down 0
## Up   1
```

```
glm.probs <- predict(glm.fits, type = "response")
glm.probs[1:10]
```

```
##      1      2      3      4      5      6      7      8
## 0.5070841 0.4814679 0.4811388 0.5152224 0.5107812 0.5069565 0.4926509 0.5092292
##      9     10
## 0.5176135 0.4888378
```

```
#Hm, it's all 50/50... de la poudre de perlinpinpin!
```

```
#Can then convert P(x/y) into labels based on their values
glm.pred <- rep("Down",1250)
glm.pred[glm.probs > 0.5] <- "Up"
#Plot predictions against training values
table(glm.pred,Direction)
```

```
##      Direction
```

```
## glm.pred Down Up
##      Down  145 141
##      Up    457 507
```

```
#Calculate Model Accuracy for training data
train_error <- (145 + 507)/1250
#Or do
mean(glm.pred == Direction)
```

```
## [1] 0.5216
```

```
## ~50% accuracy on the training data... oh boy
```

```
##Creating a test data set
train <- ( Year < 2005)
Smarket.2005 <- Smarket[!train ,]
dim (Smarket.2005)
```

```
## [1] 252  9
```

```
Direction.2005 <- Direction[!train]
```

```
#repeat log-red with data from before 2005
glm.fits1 <- glm(Direction ~ Lag1 + Lag2 + Lag3 + Lag4 + Lag5 + Volume, data =
              Smarket, family = binomial, subset = train)
glm.probs1 <- predict(glm.fits1, Smarket.2005, type = "response")
glm.pred1 <- rep("Down", 252)
glm.pred1[glm.probs1 > 0.5] <- "Up"
table(glm.pred1,Direction.2005)
```

```
##      Direction.2005
## glm.pred1 Down Up
##      Down    77 97
##      Up     34 44
```

```
mean(glm.pred1 == Direction.2005)
```

```
## [1] 0.4801587
```

```
mean(glm.pred1 != Direction.2005)
```

```
## [1] 0.5198413
```

```
#Model is worse than a guess 0o0
```

```
#Rebuild model using less parameters
#repeat log-red with data from before 2005
glm.fits2 <- glm(Direction ~ Lag1 + Lag2, data =
              Smarket, family = binomial, subset = train)
```

```
glm.probs2 <- predict(glm.fits2, Smarket.2005, type = "response")
glm.pred2 <- rep("Down", 252)
glm.pred2[glm.probs2 > 0.5] <- "Up"
table(glm.pred2, Direction.2005)
```

```
##           Direction.2005
## glm.pred2 Down  Up
##      Down   35  35
##      Up    76 106
```

```
mean(glm.pred2 == Direction.2005)
```

```
## [1] 0.5595238
```

```
# based on this confusion matrix one could develop a trading strategy
# (don't necessarily need a perfect model, just need one that improves your
# chances)
```

## Linear Discriminant Analysis

```
##LDAs are fit with the lda function
library(MASS)
```



```
##
## Attaching package: 'MASS'

## The following object is masked from 'package:ISLR2':
##
## Boston
```

```
lda.fit <- lda(Direction ~ Lag1 + Lag2, data = Smarket, subset = train)
lda.fit
```

```
## Call:
## lda(Direction ~ Lag1 + Lag2, data = Smarket, subset = train)
##
## Prior probabilities of groups:
##      Down      Up
## 0.491984 0.508016
##
## Group means:
##           Lag1      Lag2
## Down 0.04279022 0.03389409
## Up  -0.03954635 -0.03132544
##
## Coefficients of linear discriminants:
##           LD1
## Lag1 -0.6420190
## Lag2 -0.5135293
```

```
plot(lda.fit)
```



```
lda.pred <- predict(lda.fit, Smarket.2005)
names(lda.pred)
```

```
## [1] "class"      "posterior" "x"
```

```
lda.class <- lda.pred$class
table(lda.class, Direction.2005)
```

```
##           Direction.2005
## lda.class Down   Up
##      Down   35   35
##      Up    76  106
```

```
mean(lda.class == Direction.2005)
```

```
## [1] 0.5595238
```

```
#Can recreate the labels by applying a 50% threshold to the posterior
#probabilities
## May want to change posterior threshold depending on what bias we see in test
#results. Too many FP? Could try increasing the threshold
```

## Quadratic Discriminant Analysis

```
## Called with the qda() function
```

```
qda.fit <- qda(Direction ~ Lag1 + Lag2, data = Smarket, subset = train)
qda.fit
```

```
## Call:
## qda(Direction ~ Lag1 + Lag2, data = Smarket, subset = train)
##
## Prior probabilities of groups:
##      Down      Up
## 0.491984 0.508016
##
## Group means:
##           Lag1      Lag2
## Down 0.04279022 0.03389409
## Up   -0.03954635 -0.03132544
```

```
qda.class <- predict(qda.fit, Smarket.2005)$class
table(qda.class, Direction.2005)
```

```
##           Direction.2005
## qda.class Down  Up
##      Down   30  20
##      Up    81 121
```

```
#At first glance we would say that the model is good at predicting Ups However, there is a very high FP
mean(qda.class == Direction.2005)
```

```
## [1] 0.5992063
```

```
#Since the overall accuracy is 60% it's still much better than the previous models but is still a conce
```

## Naive Bayes

```
#Need to load the *e1071* library which contains the Naive Bayes method for R
library(e1071)
```

```
nb.fit <- naiveBayes(Direction ~ Lag1 + Lag2, data = Smarket, subset = train)
```

```
nb.fit
```

```
##
## Naive Bayes Classifier for Discrete Predictors
##
## Call:
## naiveBayes.default(x = X, y = Y, laplace = laplace)
```



```
##
## A-priori probabilities:
## Y
##      Down      Up
## 0.491984 0.508016
##
## Conditional probabilities:
##      Lag1
## Y      [,1]      [,2]
## Down 0.04279022 1.227446
## Up   -0.03954635 1.231668
##
##      Lag2
## Y      [,1]      [,2]
## Down 0.03389409 1.239191
## Up   -0.03132544 1.220765
```

```
nb.class <- predict(nb.fit, Smarket.2005)
table(nb.class, Direction.2005)
```

```
##      Direction.2005
## nb.class Down  Up
##      Down   28  20
##      Up    83 121
```

```
mean(nb.class == Direction.2005)
```

```
## [1] 0.5912698
```

```
#can also generate the probabilities for each prediction
nb.pred <- predict(nb.fit, Smarket.2005, type = "raw")
nb.pred[1:5,]
```

```
##      Down      Up
## [1,] 0.4873164 0.5126836
## [2,] 0.4762492 0.5237508
## [3,] 0.4653377 0.5346623
## [4,] 0.4748652 0.5251348
## [5,] 0.4901890 0.5098110
```

## K-Nearest Neighbours

```
#Performed by using the knn() function

#Need a matrix composed of predictors from training data (train.x)
#A matrix with predictors from data we are trying to predict (test.xt)
#A vector with class labels for train.x (train.Direction)
#K, the number of NN to use in classifier
library(class)
library(ISLR2)
```

```
#Setup test data  
attach(Smarket)
```

```
## The following objects are masked from Smarket (pos = 6):  
##  
##      Direction, Lag1, Lag2, Lag3, Lag4, Lag5, Today, Volume, Year
```

```
train <- ( Year < 2005)  
train.x <- cbind(Lag1,Lag2)[train,]  
test.x <- cbind(Lag1,Lag2)[!train,]  
train.Direction <- Direction[train]  
Direction.2005 <- Direction[!train]  
set.seed(1)  
knn.pred <- knn(train.x,test.x, train.Direction, k = 1)  
table(knn.pred,Direction.2005)
```

```
##      Direction.2005  
## knn.pred Down Up  
##      Down   43 58  
##      Up    68 83
```

```
#50% accuracy with k = 1, not great  
mean(knn.pred == Direction.2005)
```

```
## [1] 0.5
```

```
#K = 3  
knn.pred2 <- knn(train.x,test.x, train.Direction, k = 3)  
table(knn.pred2,Direction.2005)
```

```
##      Direction.2005  
## knn.pred2 Down Up  
##      Down   48 54  
##      Up    63 87
```

```
mean(knn.pred2 == Direction.2005)
```

```
## [1] 0.5357143
```

```
##Overall the QDA model had the highest accuracy on this data set  
#could still change what parameters we use in the model later to see if this brings us above the 60%
```

KNN on the caravan data set

```
dim(Caravan)
```

```
## [1] 5822 86
```

```
attach(Caravan)
summary(Purchase)
```

```
##      No  Yes
## 5474  348
```

```
#Scaling all the predictors to prevent any bias due to different magnitudes from affecting knn prediction
standardized.X <- scale(Caravan[, -86])
```

```
#Training subsets
```

```
test <- 1:1000
train.X1 <- standardized.X[-test,]
test.X1 <- standardized.X[test,]
train.Y1 <- Purchase[-test]
test.Y1 <- Purchase[test]

set.seed(1)
knn.pred3 <- knn(train.X1, test.X1, train.Y1, k = 1)
mean(test.Y1 != knn.pred3)
```

```
## [1] 0.118
```

```
##The error rate looks good here, but because it is higher than the proportion of Yes in the training data
mean(test.Y1 != "No")
```

```
## [1] 0.059
```

```
table(knn.pred3, test.Y1)
```

```
##           test.Y1
## knn.pred3  No  Yes
##           No  873  50
##           Yes  68   9
```

```
(9)/(66+9)
```

```
## [1] 0.12
```

```
#12% success rate if you only sell insurance to people who are predicted to buy it
# double the odds of just asking everyone (6%)
```

```
knn.pred4 <- knn(train.X1, test.X1, train.Y1, k = 3)
table(knn.pred4, test.Y1)
```

```
##           test.Y1
## knn.pred4  No  Yes
##           No  920  54
##           Yes  21   5
```

5/24

```
## [1] 0.2083333
```

*#K = 3 improves this even more!*

```
knn.pred5 <- knn(train.X1,test.X1,train.Y1, k = 5)
table(knn.pred5,test.Y1)
```

```
##           test.Y1
## knn.pred5 No Yes
##           No  930  55
##           Yes   11   4
```

4/15

```
## [1] 0.2666667
```

*#Even higher! unfortunately this comes at the cost of only suggesting 15 candidate buyers...*

*##Trying with a logistic regression*

```
glm.fits3 <- glm(Purchase ~ . , data = Caravan, subset = -test, family = binomial)
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
glm.probs3 <- predict(glm.fits3,Caravan[test,],type = "response")
glm.pred3 <- rep("No",1000)
glm.pred3[glm.probs3 > 0.5] <- "Yes"
table(glm.pred3,test.Y1)
```

```
##           test.Y1
## glm.pred3 No Yes
##           No  934  59
##           Yes   7   0
```

*#All yes predictions are wrong!*

```
glm.pred4 <- rep("No",1000)
glm.pred4[glm.probs3 > 0.25] <- "Yes"
table(glm.pred4,test.Y1)
```

```
##           test.Y1
## glm.pred4 No Yes
##           No  919  48
##           Yes  22  11
```

*#Now have a 33% success rate which is the best we have seen so far*

## Poisson Regression

```
attach(Bikeshare)
dim(Bikeshare)
```

```
## [1] 8645 15
```

```
names(Bikeshare)
```

```
## [1] "season"      "mnth"        "day"         "hr"          "holiday"
## [6] "weekday"     "workingday"  "weathersit"   "temp"        "atemp"
## [11] "hum"         "windspeed"   "casual"      "registered"  "bikers"
```

```
#Testing linear regression
```

```
mod.lm <- lm(bikers ~ mnth + hr + workingday + temp + weathersit, data = Bikeshare)
summary(mod.lm)
```

```
##
## Call:
## lm(formula = bikers ~ mnth + hr + workingday + temp + weathersit,
##     data = Bikeshare)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -299.00  -45.70   -6.23   41.08  425.29
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    -68.632     5.307  -12.932 < 2e-16 ***
## mnthFeb           6.845     4.287    1.597 0.110398
## mnthMarch        16.551     4.301    3.848 0.000120 ***
## mnthApril        41.425     4.972    8.331 < 2e-16 ***
## mnthMay          72.557     5.641   12.862 < 2e-16 ***
## mnthJune         67.819     6.544   10.364 < 2e-16 ***
## mnthJuly         45.324     7.081    6.401 1.63e-10 ***
## mnthAug          53.243     6.640    8.019 1.21e-15 ***
## mnthSept         66.678     5.925   11.254 < 2e-16 ***
## mnthOct          75.834     4.950   15.319 < 2e-16 ***
## mnthNov          60.310     4.610   13.083 < 2e-16 ***
## mnthDec          46.458     4.271   10.878 < 2e-16 ***
## hr1             -14.579     5.699   -2.558 0.010536 *
## hr2             -21.579     5.733   -3.764 0.000168 ***
## hr3             -31.141     5.778   -5.389 7.26e-08 ***
## hr4             -36.908     5.802   -6.361 2.11e-10 ***
## hr5             -24.135     5.737   -4.207 2.61e-05 ***
## hr6              20.600     5.704    3.612 0.000306 ***
## hr7            120.093     5.693   21.095 < 2e-16 ***
## hr8            223.662     5.690   39.310 < 2e-16 ***
## hr9            120.582     5.693   21.182 < 2e-16 ***
## hr10             83.801     5.705   14.689 < 2e-16 ***
## hr11            105.423     5.722   18.424 < 2e-16 ***
## hr12            137.284     5.740   23.916 < 2e-16 ***
## hr13            136.036     5.760   23.617 < 2e-16 ***
```

```
## hr14                126.636      5.776  21.923 < 2e-16 ***
## hr15                132.087      5.780  22.852 < 2e-16 ***
## hr16                178.521      5.772  30.927 < 2e-16 ***
## hr17                296.267      5.749  51.537 < 2e-16 ***
## hr18                269.441      5.736  46.976 < 2e-16 ***
## hr19                186.256      5.714  32.596 < 2e-16 ***
## hr20                125.549      5.704  22.012 < 2e-16 ***
## hr21                 87.554      5.693  15.378 < 2e-16 ***
## hr22                 59.123      5.689  10.392 < 2e-16 ***
## hr23                 26.838      5.688   4.719 2.41e-06 ***
## workingday           1.270      1.784   0.711 0.476810
## temp                157.209     10.261  15.321 < 2e-16 ***
## weathersitcloudy/misty -12.890      1.964  -6.562 5.60e-11 ***
## weathersitlight rain/snow -66.494      2.965 -22.425 < 2e-16 ***
## weathersitheavy rain/snow -109.745     76.667  -1.431 0.152341
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 76.5 on 8605 degrees of freedom
## Multiple R-squared:  0.6745, Adjusted R-squared:  0.6731
## F-statistic: 457.3 on 39 and 8605 DF, p-value: < 2.2e-16
```

```
#Changing data values for mnth and hr
contrasts(Bikeshare$hr) <- contr.sum(24)
contrasts(Bikeshare$mnth) <- contr.sum(12)
mod.lm2 <- lm(bikers ~ mnth + hr + workingday + temp + weathersit, data = Bikeshare)
summary(mod.lm2)
```


```
##
## Call:
## lm(formula = bikers ~ mnth + hr + workingday + temp + weathersit,
##     data = Bikeshare)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -299.00  -45.70   -6.23   41.08  425.29
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    73.5974     5.1322  14.340 < 2e-16 ***
## mnth1         -46.0871     4.0855 -11.281 < 2e-16 ***
## mnth2         -39.2419     3.5391 -11.088 < 2e-16 ***
## mnth3         -29.5357     3.1552  -9.361 < 2e-16 ***
## mnth4          -4.6622     2.7406  -1.701  0.08895 .
## mnth5          26.4700     2.8508   9.285 < 2e-16 ***
## mnth6          21.7317     3.4651   6.272 3.75e-10 ***
## mnth7          -0.7626     3.9084  -0.195  0.84530
## mnth8           7.1560     3.5347   2.024  0.04295 *
## mnth9          20.5912     3.0456   6.761 1.46e-11 ***
## mnth10         29.7472     2.6995  11.019 < 2e-16 ***
## mnth11         14.2229     2.8604   4.972 6.74e-07 ***
## hr1           -96.1420     3.9554 -24.307 < 2e-16 ***
## hr2          -110.7213     3.9662 -27.916 < 2e-16 ***
## hr3          -117.7212     4.0165 -29.310 < 2e-16 ***
```

```
## hr4          -127.2828      4.0808 -31.191 < 2e-16 ***
## hr5          -133.0495      4.1168 -32.319 < 2e-16 ***
## hr6          -120.2775      4.0370 -29.794 < 2e-16 ***
## hr7          -75.5424       3.9916 -18.925 < 2e-16 ***
## hr8           23.9511       3.9686   6.035 1.65e-09 ***
## hr9          127.5199       3.9500  32.284 < 2e-16 ***
## hr10         24.4399       3.9360   6.209 5.57e-10 ***
## hr11        -12.3407       3.9361  -3.135 0.00172 **
## hr12          9.2814       3.9447   2.353 0.01865 *
## hr13         41.1417       3.9571  10.397 < 2e-16 ***
## hr14         39.8939       3.9750  10.036 < 2e-16 ***
## hr15         30.4940       3.9910   7.641 2.39e-14 ***
## hr16         35.9445       3.9949   8.998 < 2e-16 ***
## hr17         82.3786       3.9883  20.655 < 2e-16 ***
## hr18        200.1249       3.9638  50.488 < 2e-16 ***
## hr19        173.2989       3.9561  43.806 < 2e-16 ***
## hr20         90.1138       3.9400  22.872 < 2e-16 ***
## hr21         29.4071       3.9362   7.471 8.74e-14 ***
## hr22         -8.5883       3.9332  -2.184 0.02902 *
## hr23        -37.0194       3.9344  -9.409 < 2e-16 ***
## workingday     1.2696       1.7845   0.711 0.47681
## temp        157.2094      10.2612  15.321 < 2e-16 ***
## weathersitcloudy/misty -12.8903   1.9643 -6.562 5.60e-11 ***
## weathersitlight rain/snow -66.4944   2.9652 -22.425 < 2e-16 ***
## weathersitheavy rain/snow -109.7446  76.6674 -1.431 0.15234
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 76.5 on 8605 degrees of freedom
## Multiple R-squared:  0.6745, Adjusted R-squared:  0.6731
## F-statistic: 457.3 on 39 and 8605 DF, p-value: < 2.2e-16
```

```
## Showing that both coding approaches do not change the model's predictions
all.equal(predict(mod.lm), predict(mod.lm2))
```

```
## [1] TRUE
```

```
## Get month coefficients
coef.months <- c(coef(mod.lm2)[2:12], -sum(coef(mod.lm2)[2:12]))
plot(coef.months, xlab = "Month", ylab = "Coefficient", xaxt = "n", col = "blue", pch = 19, type = "o")
axis(side = 1, at = 1:12, labels = c("J", "F", "M", "A", "M", "J", "J", "A", "S", "O", "N", "D"))
```



```
## now fitting a poisson regression on the data instead
```

```
mod.pois <- glm(bikers ~ mnth + hr + workingday + temp + weathersit, data = Bikeshare, family = poisson)
summary(mod.pois)
```

```
##
## Call:
## glm(formula = bikers ~ mnth + hr + workingday + temp + weathersit,
##      family = poisson, data = Bikeshare)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -20.7574  -3.3441  -0.6549   2.6999  21.9628
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)      4.118245   0.006021  683.964 < 2e-16 ***
## mnth1           -0.670170   0.005907 -113.445 < 2e-16 ***
## mnth2           -0.444124   0.004860 -91.379 < 2e-16 ***
## mnth3           -0.293733   0.004144 -70.886 < 2e-16 ***
## mnth4            0.021523   0.003125   6.888 5.66e-12 ***
## mnth5            0.240471   0.002916  82.462 < 2e-16 ***
## mnth6            0.223235   0.003554  62.818 < 2e-16 ***
## mnth7            0.103617   0.004125  25.121 < 2e-16 ***
## mnth8            0.151171   0.003662  41.281 < 2e-16 ***
## mnth9            0.233493   0.003102  75.281 < 2e-16 ***
```



```


## mnth10          0.267573    0.002785    96.091 < 2e-16 ***
## mnth11          0.150264    0.003180    47.248 < 2e-16 ***
## hr1            -0.754386    0.007879   -95.744 < 2e-16 ***
## hr2            -1.225979    0.009953  -123.173 < 2e-16 ***
## hr3            -1.563147    0.011869  -131.702 < 2e-16 ***
## hr4            -2.198304    0.016424  -133.846 < 2e-16 ***
## hr5            -2.830484    0.022538  -125.586 < 2e-16 ***
## hr6            -1.814657    0.013464  -134.775 < 2e-16 ***
## hr7            -0.429888    0.006896   -62.341 < 2e-16 ***
## hr8             0.575181    0.004406   130.544 < 2e-16 ***
## hr9             1.076927    0.003563   302.220 < 2e-16 ***
## hr10            0.581769    0.004286   135.727 < 2e-16 ***
## hr11            0.336852    0.004720    71.372 < 2e-16 ***
## hr12            0.494121    0.004392   112.494 < 2e-16 ***
## hr13            0.679642    0.004069   167.040 < 2e-16 ***
## hr14            0.673565    0.004089   164.722 < 2e-16 ***
## hr15            0.624910    0.004178   149.570 < 2e-16 ***
## hr16            0.653763    0.004132   158.205 < 2e-16 ***
## hr17            0.874301    0.003784   231.040 < 2e-16 ***
## hr18            1.294635    0.003254   397.848 < 2e-16 ***
## hr19            1.212281    0.003321   365.084 < 2e-16 ***
## hr20            0.914022    0.003700   247.065 < 2e-16 ***
## hr21            0.616201    0.004191   147.045 < 2e-16 ***
## hr22            0.364181    0.004659    78.173 < 2e-16 ***
## hr23            0.117493    0.005225    22.488 < 2e-16 ***
## workingday      0.014665    0.001955    7.502 6.27e-14 ***
## temp            0.785292    0.011475    68.434 < 2e-16 ***
## weathersitcloudy/misty -0.075231  0.002179   -34.528 < 2e-16 ***
## weathersitlight rain/snow -0.575800  0.004058  -141.905 < 2e-16 ***
## weathersitheavy rain/snow -0.926287  0.166782   -5.554 2.79e-08 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for poisson family taken to be 1)
##
## Null deviance: 1052921 on 8644 degrees of freedom
## Residual deviance: 228041 on 8605 degrees of freedom
## AIC: 281159
##
## Number of Fisher Scoring iterations: 5

```


```

coef.months2 <- c(coef(mod.pois)[2:12], -sum(coef(mod.pois[2:12])))
plot(coef.months2, xlab = "Month", ylab = "Coefficient", xaxt = "n", col = "blue", pch = 19, type = "o",
axis(side = 1, at = 1:12, labels = c("J", "F", "M", "A", "M", "J", "J", "A", "S", "O", "N", "D")))


```



```
coef.hr <- c(coef(mod.pois)[13:35], -sum(coef(mod.pois[13:35])))  
plot(coef.hr, xlab = "Hour", ylab = "Coefficient", col = "red", pch = 18, type = "o")
```



```
plot(predict(mod.lm2),predict(mod.pois, type = "response"))  
abline(0,1,col = 2, lwd = 3)
```



Lab is complete